Recap: preserve meaning and provenance

Explicit vocabulary, qualified statements and inspectable evidence

1. Guide extraction explicitly

Model and prompt choices produce different graph views.
Ontology guidance still needs validation.

2. Preserve each claim’s conditions

Keep period, scope, unit and source with the statement.
Manage competing builds through qualified or separate database views.

3. Evaluate the downstream answer

317 questions: higher 66 · equal 193 · lower 58.
The selected ontology condition produces mixed answer-quality results.

459 nonempty builds · 1,156 entity nodes · 4,604 admitted statements

Namespace-qualified sums across two extraction conditions; not canonical world entities.

Same-model automated grading; 0 independent human reviews. More stored statements do not establish more correct facts.

Recap: engineer evidence delivery

Overlap independent reads, then inspect both median and tail latency

1. Assemble evidence before generation

Independent reads → one complete tool message → one model answer.
Preserve the same evidence when comparing execution paths.

2. Overlap waiting when branches are independent

Sequential → parallel median, concurrency 1:
Python 3.127 → 2.074 ms · Rust 2.094 → 1.303 ms.

3. Inspect the tail with the median

Concurrency 4: Rust / Python p50 = 3.433 / 7.668 ms;
Rust / Python p95 = 16.832 / 14.613 ms.

Similar PROFILE work: 133 vs 140 DB hits · same 14 statements

Measure evidence collection cost, payload parity and p50/p95 together.

DB/client timings exclude model generation. GIL causality, native Composite effects and intra-query runtime were not isolated.

XCENA - Graph Product Engineer - Yitae Jeong

Thank you

Code, aggregate results and reproducibility notes

---

github.com/tteon/nodes26

Questions & discussion · Session track: Data Intelligence